

Computer Networks and Distributed Systems

Tutorial – Bridging and Routing Protocols

Course 527 – Spring Term 2015-2016

Emil C Lupu and Daniele Sgandurra

e.c.lupu@imperial.ac.uk, d.sgandurra@imperial.ac.uk

Part 4 – Contents

Spanning Tree Bridging Protocol

Distance Vector Routing Protocol

Link-State Routing Protocol

1

Spanning Tree Protocol (Layer 2)

Used by switches to turn a redundant topology into a **spanning tree**:

- Avoid and **eliminate loops** in the network by negotiating a loop-free path through a **root bridge**.
- Disables unwanted links by **blocking** ports.
- Ensures that there will be **only one active path** to every destination.

STP defined by **IEEE 802.1d**.

Switches run STP by default – no configuration needed.

2

Spanning Tree Algorithm

The switches use this algorithm to decide which ports should be shut down.

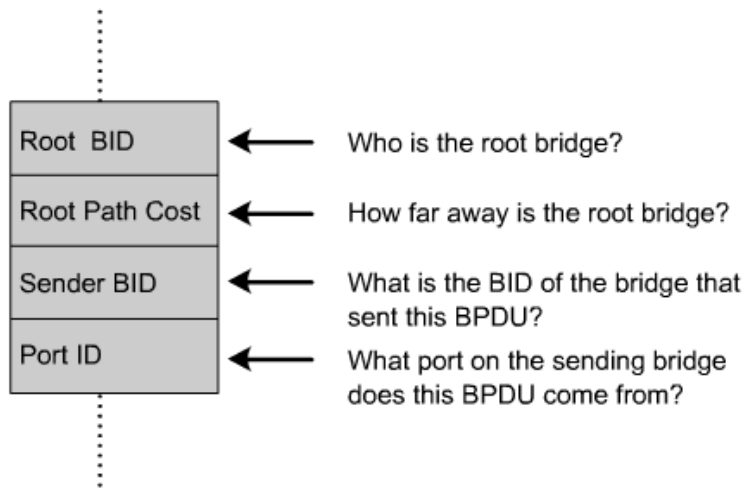
1. Choose one switch to be “**root bridge**”
2. Choose a “**root port**” on each other switch
3. Choose a “**designated port**” on each segment.
4. **Close down** all other ports.

See step-by-step example at:

http://www.cisco.com/image/gif/paws/10556/spanning_tree1.swf

3

BPDU (Bridge Protocol Data Units)



4

Root and Designated Ports

Root Port: This is the port that is the closest to the root bridge in terms of path cost.

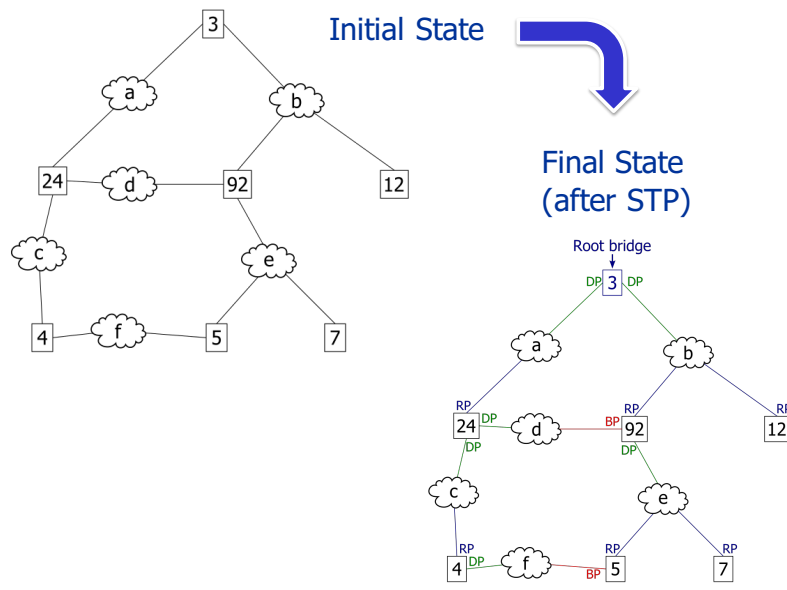
Designated port: it can send the best BPDU on the segment to which it is connected.

- On a **given segment**, there can be **only one path toward the root bridge** otherwise redundant paths would create a bridging loop.
- All bridges connected to a given segment listen to each other's BPDUs and agree on the bridge sending the best BPDU as the designated bridge for the segment.

A Root Port can **never** be a Designated Port.

5

Spanning Tree Example



6

Distance Vector Protocol (Layer 3)

A **distance vector routing** algorithm operates by having each router maintain a table (i.e., vector):

- **Best known distance to each router.**
- **Which link to use to get there.**

Each table is indexed by, and containing one entry for each router in the network:

- Preferred outgoing line for that destination.
- Estimate of the distance to that destination.

7

Distance Vector Protocol

A router is assumed to know the **distance** (cost) to each of its neighbors.

Once every T_{msec} , each router sends to each neighbor a list of its **estimated delays** to each destination:

- It also receives a similar list from each neighbor.

If R has neighbor X:

- Table from X: X_i = estimated delay to get to router i.
- If delay to X estimated by R is m:
 - R can reach i in $m + X_i$ msec.

8

Bellman-Ford Algorithm

$c(x, v)$ = cost for direct link from x to v :

- Node x maintains costs of direct links $c(x, v)$.

$D_x(y)$ = estimate of least cost from x to y :

- Node x maintains distance vector $D_x = [D_x(y): y \in N]$.

Node x maintains its neighbors' distance vectors:


- For each neighbor v , x maintains $D_v = [D_v(y): y \in N]$.

Each node v periodically sends D_v to its neighbors:

- And neighbors update their own distance vectors.
- $D_x(y) \leftarrow \min_v \{c(x, v) + D_v(y)\}$ for each node $y \in N$.

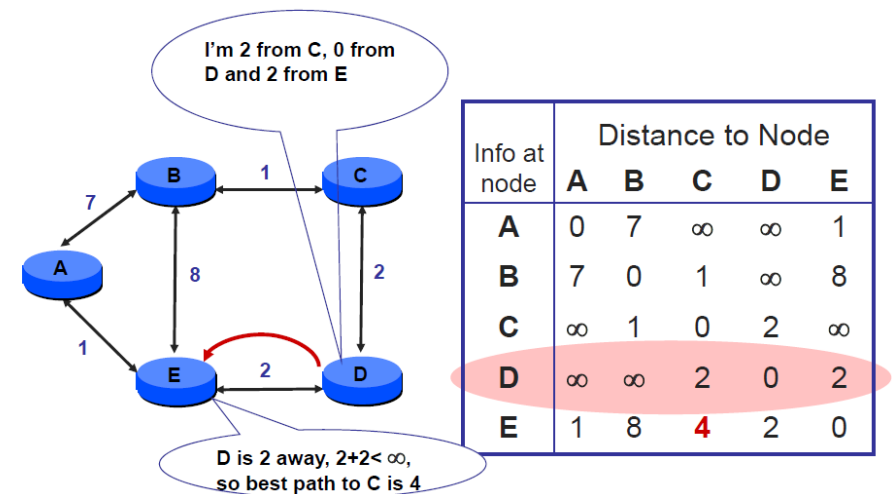
9

Example: Initial State



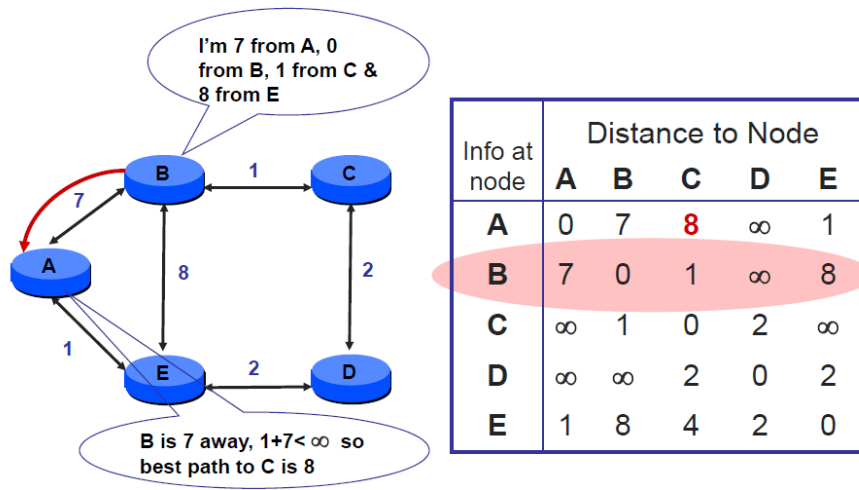
10

D Sends Vector to E



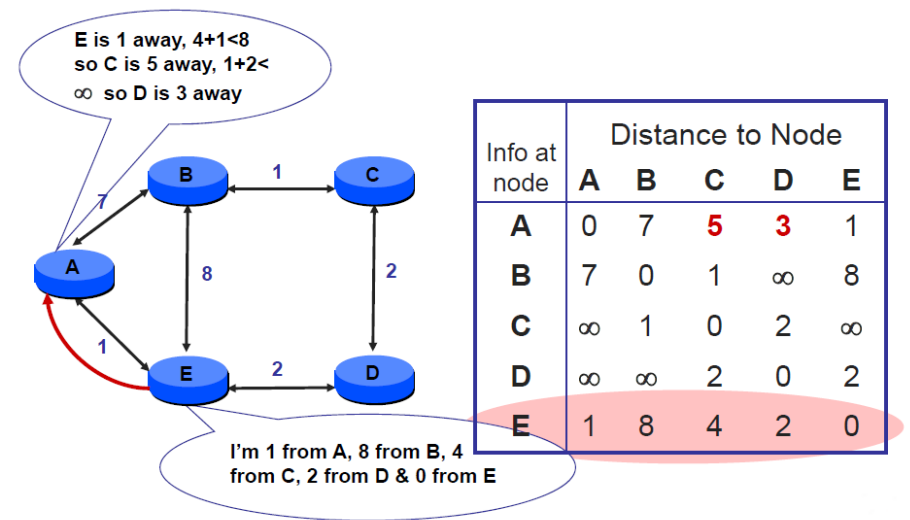
11

B Sends Vector to A




12

E Sends Vector to A




13

Until Convergence




14

Node B's Distance Vector



15

Counting to Infinity Problem



16

Link-State Routing (Layer 3)

Same goal as Distance Vector, but different approach:

- Tell **every** node what you know about your **direct neighbors**.

Two phases:

- **Reliable flooding:**
 - Tell all routers what you know about your local topology.
- **Path calculation** (Dijkstra's algorithm):
 - Each router computes best path over complete network.

17

Reliable-Flooding

Each router transmits a **Link State Packet** on all links:

- Identifier of the sender.
- List of neighbors with their associated costs.

A neighboring router **forwards** out all links except incoming.

Keep a copy locally; don't forward previously-seen LSPs.

Caveats:

- Acknowledgments and retransmissions.
- Sequence numbers.
- Time-to-live for each packet.

18

Running Dijkstra Algorithm

Determine the **shortest distance** between a start node and any other node in a graph.


```

1: function Dijkstra(Graph, source):
2:   for each vertex v in Graph:      // Initialization
3:     dist[v] := infinity           // Initial distance from source to v is set to ∞
4:     previous[v] := undefined      // Previous node in optimal path from source
5:   dist[source] := 0                // Distance from source to source
6:   Q := the set of all nodes in Graph // All nodes in the graph are unoptimized
7:   while Q is not empty:           // main loop
8:     u := node in Q with smallest dist[ ]
9:     remove u from Q
10:    for each neighbor v of u:      // where v has not yet been removed from Q.
11:      alt := dist[u] + dist_between(u, v)
12:      if alt < dist[v] // Relax (u,v)
13:        dist[v] := alt
14:        previous[v] := u
15:  return previous[ ]
  
```

19

Dijkstra's LSR Algorithm

Consider the following network:




Link state database:

A	B	C	D	E	F	G
B 6	A 6	B 6	A 2	B 1	C 2	C 5
D 2	C 2	F 2	E 2	D 2	E 4	F 1
	E 1	G 5		F 4	G 1	



Dijkstra's LSR Algorithm

Running Dijkstra on C





Dijkstra's LSR Algorithm

Examine C's LSP
– Cost to F, G, and B





Dijkstra's LSR Algorithm

Pick node in Q with lowest dist: F
– Its predecessor is determined: C (cannot be changed)





Dijkstra's LSR Algorithm

Update G cost and predecessor





Dijkstra's LSR Algorithm

Pick node in Q with lowest dist: B





Dijkstra's LSR Algorithm

Update E cost and predecessor




Dijkstra's LSR Algorithm

Pick node in Q with lowest dist: E




Dijkstra's LSR Algorithm

Pick node in Q with lowest dist: G




Dijkstra's LSR Algorithm

Pick node in Q with lowest dist: D
– Update A cost and predecessor




Dijkstra's LSR Algorithm

Pick node in Q with lowest dist: A



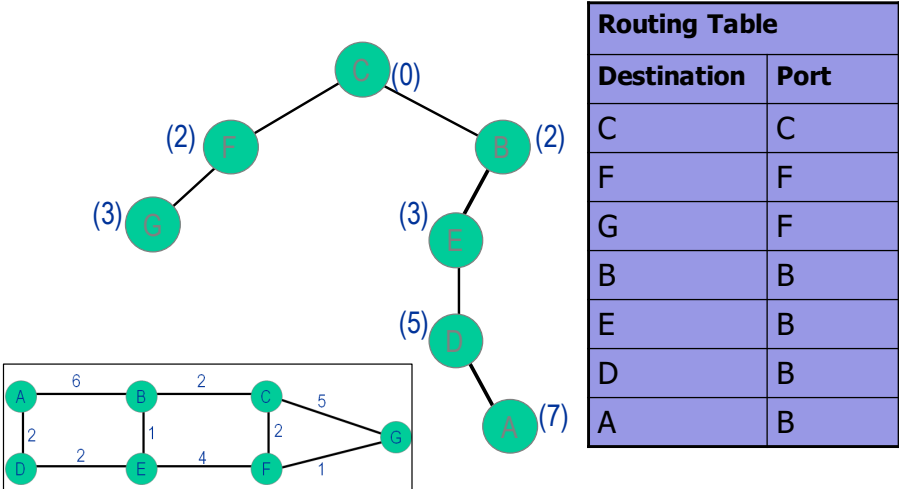
Dijkstra's LSR Algorithm

Q is empty: DONE



Dijkstra's LSR Algorithm

We can now create a routing (forwarding) table for C:



Routing Table	
Destination	Port
C	C
F	F
G	F
B	B
E	B
D	B
A	B