

## Computer Networks and Distributed Systems

### Part 5 – Transport Layer

Course 527 – Spring Term 2015-2016

**Emil C Lupu & Daniele Sgandurra**

[e.c.lupu@imperial.ac.uk](mailto:e.c.lupu@imperial.ac.uk), [d.sgandurra@imperial.ac.uk](mailto:d.sgandurra@imperial.ac.uk)

## Host-to-Host Communications

Datagrams transferred between hosts

- Network/Internet level in stack

Routed through networks based on IP addresses

- Source address of sending machine
- Destination address of recipient machine

But:

- No identification of which applications send or receive
- Only unreliable connection-less datagram service

## Part 5 – Contents

Transport Layer

- End-to-end communications
- End-to-end addressing: ports

Transport protocols

- **UDP**
  - Header format
- **TCP**
  - Header format
  - Connection setup
  - Retransmission
  - Flow / congestion control
  - ...

1

## App-to-App Communications

Identify application instances (processes)?

- Need well-understood identification
- Need to handle process restart
- Need to handle pools of processes serving clients
  - e.g. multi-threaded web server

Destination given by function or service

- One process may handle multiple services

## Protocol Ports

Use **ports** to define end points

- Abstract addressing
- 16 bit unsigned integer: 0 - 65535

Processes specify **ports** to send to or receive from

- Use operating system calls to **bind** to port
- Packets have source and destination ports

4

## Well Known Service Ports

Other systems know which port to address to

- Standard services usually run on well known ports

List of well known services (RFC 1060)

Ports 0 – 255: Internet Assigned Numbers Authority (IANA)

Ports 0 – 1023: Privileged UNIX standard services

Static file with service/port mappings

Unix: `/etc/services`

Windows: `\windows\system32\drivers\etc\services`

5

## `/etc/services`

```
tcpmux      1/tcp      # TCP port service multiplexer
echo        7/tcp
echo        7/udp
discard     9/tcp      sink null
discard     9/udp      sink null
systat      11/tcp      users
daytime     13/tcp
daytime     13/udp
netstat     15/tcp
qotd        17/tcp      quote
msp         18/tcp      # message send protocol
msp         18/udp
chargen     19/tcp      ttytst source
chargen     19/udp      ttytst source
ftp-data    20/tcp
ftp         21/tcp
fsp         21/udp      fspd
ssh         22/tcp      # SSH Remote Login Protocol
ssh         22/udp
telnet      23/tcp
smtp        25/tcp      mail
```

6

## Less Well Known Addresses

Ports need not all be “well known”

Source port can be generated by application

- Useful when just needed for interaction

Destination port need not be well known

- Pass information to sender for novel application

Portmapper service (e.g., for Remote Procedure Calls)

- Process register (name, port) binding
- Client can query portmapper service by name.
- Portmapper (port 111)
- See remote procedure calls later on in the course

7

## UDP/TCP Services

Two main transport layer protocols: **UDP** and **TCP**

Some services use UDP:

- **bootp** (on Windows), **tftp**, **snmp**, ...

Some services use TCP:

- **smtp**, **ftp**, **telnet**, **finger**, **http**, ...

Some services use either:

- **echo**, **dns**, **ssh**, **bootp** (on Linux), **irc**, ...

8

## Sockets

**Socket:** Abstract communications endpoint

- **Stream socket** (SOCK\_STREAM) which uses TCP
- **Datagram socket** (SOCK\_DGRAM) which uses UDP

**Active socket**

- Connected to remote active socket via open connection
- Closing connection closes sockets at both ends

**Passive socket**

- Unconnected socket awaiting incoming connection
- Active socket spawned to handle connection
  - Passive socket goes back to waiting for new connections

10

## Complete Addresses

The complete addressing for a datagram (UDP/TCP) describes the sender and receiver in such a way that the services which are communicating are fully defined:

Source IP address and source port

- Source port = return addr (may be omitted in UDP)
- e.g. `columbia.doc.ic.ac.uk:57992`

Destination IP address and destination port

- e.g. `www.amazon.co.uk:80`


9

## Sockets and Ports

A socket is not a port (but they are related)

1+ sockets associated with 1 port

- At most one passive socket
  - Awaiting new connections
- Multiple active sockets
  - Handling open connections




Hence 4-part addressing

i.e. (src IP, src port, dst IP, dst port)

- Many connections use same port
- Need to know where connection is coming from

11

# User Datagram Protocol (UDP)




## UDP provides plain, IP-like service

- Connection-less datagrams, unreliable delivery, no sequence control, possible duplication
- Good for fast transfer with resilience to packet loss

12

# UDP Header Format



## Source port

- Optional, 0 if not used, reply-to port if used

## Destination port

- Host addressing still provided by IP headers

## Length (in bytes)

- Includes header and data (min. 8 for no data)

13

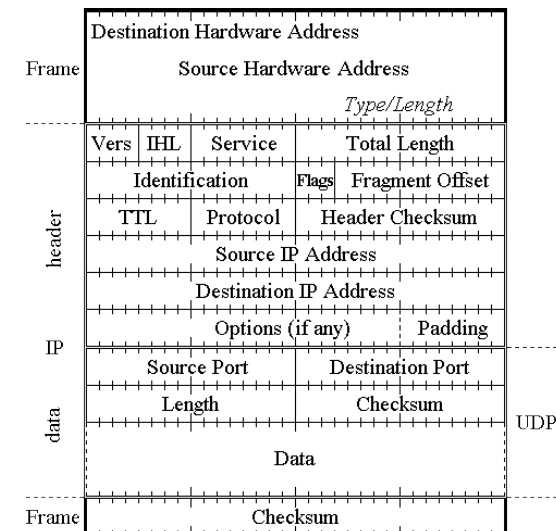
# UDP Checksum

## Checksum (16 bit using 1s complement sum)

- Calculated over pseudo header + UDP header + data
- **Pseudo header** mimics IP header
  - Source IP address
  - Destination IP address
  - Protocol (17)
  - UDP length
    - Zeros to pad to multiple of two octets
- Allows detection of changed IP headers by gateways without actually duplicating data

14

# UDP/IP Packet in Ethernet Frame



15

## Reliable Service?

UDP is unreliable

Features of reliable a service


- Unstructured stream abstraction
- Virtual circuit connection
- Full-duplex connection

Could build reliable service over UDP

- Needs to keep track of successfully transmitted packets
- Add error-correcting & retransmission mechanisms

16

## Transmission Control Protocol (TCP)



TCP adds a lot to IP:

- Streams with reliable delivery
- Full-duplex operation
- Flow control
- Network adaptation for congestion control
- Complexity & overheads

17

## Connections

TCP uses **connection** as its basic abstraction

- Rather than just protocol port (receiving datagrams)

Connection-oriented service on top of connectionless IP

- Connection identified by pair of endpoints

**Endpoint** is (host, port) tuple: (IP addr:port)

e.g. src: **146.169.7.41:1069**  
dst: **140.247.60.24:25**

Mail application on **146.169.7.41** connects to  
SMTP port on **140.247.60.24**

18

## TCP Features

**Streams**

- TCP data is stream of bytes
- Underlying datagrams concealed

**Sequence numbers** for reliable delivery

- Used to maintain byte order in stream
- TCP detects lost data and arranges retransmission
- Stream delayed during retransmission to maintain byte sequence

19

### Flow control

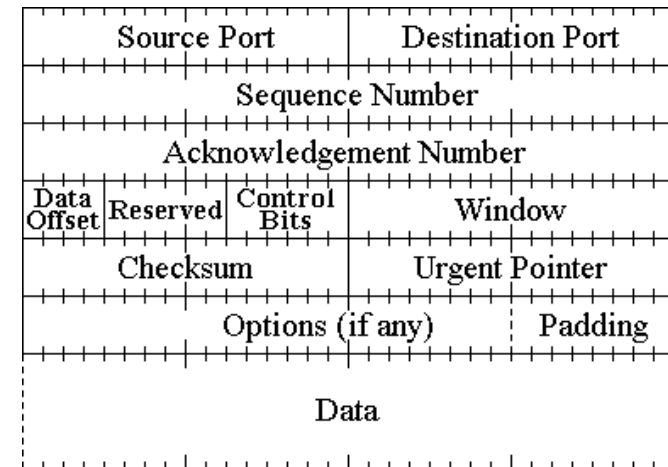
- Manages data buffers and coordinates traffic to prevent overflows
- Fast senders have to pause for slow receivers to keep up

### Congestion control

- Monitors and learns delay characteristics of network
- Adjusts operation to maximise throughput without overloading network

20

## TCP Segment Format



21

## TCP Fields: Sequence Numbers

### Sequence number (seq num)

- Indicates position in stream of 1<sup>st</sup> data byte in segment

### Acknowledgement number (ack num)

- Exploit full-duplex connection and use segments to piggyback acknowledgments
- Ack num = next seq num sender expects to receive

22

## TCP Fields: Data Sizes

### Data offset: Number of 32-bit words in TCP header

- Needed because of variable length options

### Window size (used for flow control)

- Num of data bytes which may be sent, starting with byte indicated by ack num
- Recipient should not send more than (window size – bytes sent) without acks (in transit)
- 0 means “no more data now, please”

23

## TCP Fields: Control + Checksum

### Control bits

**URG:** urgent pointer valid

**ACK:** ack num valid

**PSH:** “push” this segment  
(transmit promptly)

**RST:** reset connection

**SYN:** synchronise sequence  
numbers

**FIN:** sender has reached  
end of byte stream

### Urgent pointer

- Pointer to high priority data in stream (e.g. error conditions)

### Options

- Negotiate max segment size, scaling factor for window size, ...

### Checksum (16 bit using 1s complement sum)

- Same as for UDP with pseudo header

24

## Passive and Active Opens

Both endpoints cooperate to open TCP connection

### **Passive open**

- One end waits for incoming requests (server)

### **Active open**

- Other end initiates communication (client)

25

## Connection Control

### Two hosts must synchronise seq nums

- Controls packet order and detects loss + duplicates

### Use SYN segments to establish connection

- Establish initial sequence num (**ISN**)
- Stream positions are offsets from ISN
  - 1<sup>st</sup> data byte in segment = ISN + 1

### ISN chosen randomly

- Need to be unique over life-time of connection
- Starting at 0 bad idea because of old packets...


26

## Connection Establishment

### **Three way handshake**


- Establishes connection
- Sender and receiver agree on seq nums
- Works when two hosts establish connection concurrently

What happens when an old duplicate of the first message arrives?



27

## Asymmetric Connection Release




28

## Symmetric Connection Release

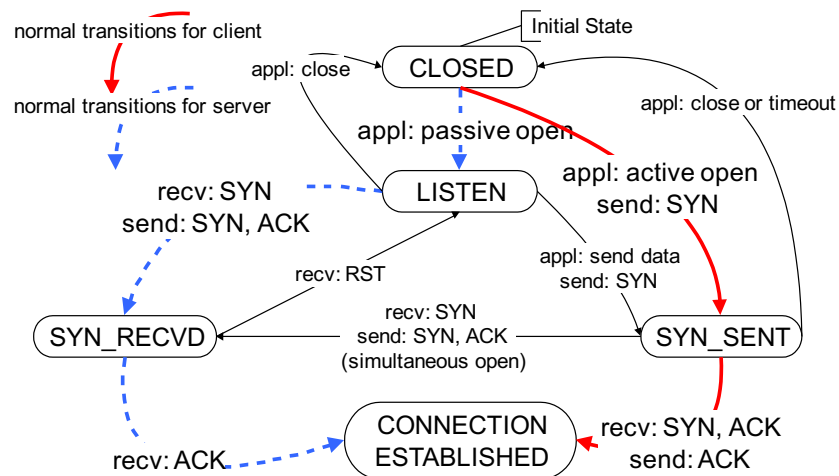
Treat connection as two unidirectional connections to be released

- Hosts agree on end seq nums
- Timeouts to handle lost messages



29


## TCP State Chart



30

## TCP Window Management

- Host 2 has 4k buffer
- Sent data controlled by WINDOW field
- Sender does not have to fill receiver's buffer with each segment
- SEQUENCE and ACKNOWLEDGEMENT indicate what has been sent/received




31



## Loss & Duplicate ACKs

### Duplicate ACK


- Generated when out-of-order segment received
- Notifies sender of duplicate and expected seq num
- Sender resends data if ACK takes longer than timeout or several duplicates received



32

## Congestion

Congestion occurs in routers and in medium access



- Multiple incoming links can saturate single outgoing link
- Slower outgoing link can be saturated by one incoming link

Routers use store-forward

- Process each packet before sending
- If buffer becomes full, drops packets

33

## TCP Congestion Control

Packet loss mostly due to congestion and not error

- Detect congestion by considering packet loss
- Change transmission rate to adapt to congestion

TCP sender maintains 2 windows

- **Receiver window** (flow control) and **congestion window**
- Uses whichever currently smaller

34

## TCP Congestion Window

Congestion window based on network conditions

- Windows grows and shrinks based on packet loss
- Different algorithms for finding optimal size e.g. additive increase, multiplicative decrease

Requires efficient timeouts to detect loss


- TCP measures RTT and adjusts timeouts

Lots of complexity to ensure throughput, fairness, ...

35

## Rough Behaviour

- Start with small window size.
- For each data segment acknowledged before time out increase window size by another segment until threshold.
- Increase linearly afterwards.
- For each packet lost reduce threshold by half.



36

37

## Summary: UDP or TCP?

### UDP

No need for reliability and error detection

- Message exchanges without transactional behaviour, e.g. DNS, DHCP
- Real-time apps, e.g. sensor monitoring, video streaming

Good for short communications

Efficient for fast networks

### TCP

Need for reliability, error correction, flow and congestion control, or security

- Terminal sessions, e.g. SSH, Telnet
- Large data transfer, e.g. web, FTP, email

Efficient for long-lived connections

Requires more CPU time and bandwidth than UDP

38